

Easy data manipulation

`{tidyr}` // `{dplyr}`

`{tidyr}`

Tidying data.

Why should we tidy data?


- Consistent data structure makes your life easier!
- 'Tidy data' helps you to:
 - Aggregate and explore (`dplyr`)
 - Visualise (`ggplot2`)
 - Model (`lm...`)
 - Re-use previous code

What is tidy data?

- Tidy data consists of **key/value** pairs
 - Key/value pairs are associations between variable names and observations of that variable
 - A key is a variable name (e.g. CITY)
 - A value is an observation of that variable (e.g. 'Glasgow')
- Columns: variables
- Rows: observations

What is tidy data?

- What does this data set **mean**?
- What are its 3 defining variables?



	Dark morph	Light morph
Male	8	2
Female	4	6

What is tidy data?

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

Mapping meaning to structure

	Dark morph	Light morph
Male	8	2
Female	4	6

→

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

Mapping meaning to structure

Each column is a variable.

Each row is an observation.

A row shows the count for one morph, for one sex.

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

	Dark	Light
Male	8	2
Female	4	6

→

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

- How do we gather the first table into a more meaningful structure?
 - Data
 - New variable names
 - The columns we want to gather

	Dark	Light
Male	8	2
Female	4	6

→

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

- How do we gather the first table into a more meaningful structure?
 - Data
 - New variable names
 - The columns we want to gather

	Dark	Light
Male	8	2
Female	4	6

→

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

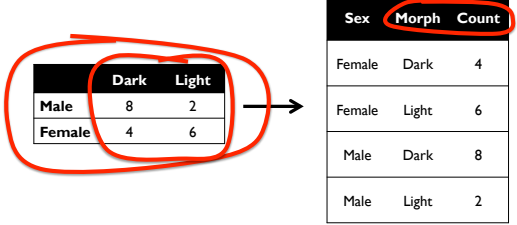
- How do we gather the first table into a more meaningful structure?
 - Data
 - New variable names
 - The columns we want to gather

	Dark	Light
Male	8	2
Female	4	6

→

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

- How do we gather the first table into a more meaningful structure?
 - Data
 - New variable names
 - The columns we want to gather



```
gather(MOTHS,
       Morph,
       Count,
       Dark:Light)
# data
# new key
# new key
# to gather
```

{dplyr}

Manipulating data.

Why manipulate data?

- Data manipulation skills make it easier to:
 - Explore and rearrange your data
 - Create summary statistics
 - Find and filter outliers
 - Create new variables based on existing observations
 - Look cool in front of all your friends

Data manipulation

- Figure out what you want to do...
- Describe this precisely in R code...
- Execute the code.
- **dplyr** makes this fast and easy.



dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

```
filter(Morph == "Dark")
```

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

```
filter(Morph == "Dark",  
       Sex == "Female")
```

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

```
arrange(desc(Count))
```

Sex	Morph	Count
Male	Dark	8
Female	Light	6
Female	Dark	4
Male	Light	2

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

```
select(Sex, Count)
```

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

```
select(-Morph)
```

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

```
rename(N = Count)
```

Sex	Morph	N
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

distinct(Morph)

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

Split-Apply-Combine

1. Split up your original data
2. Apply a function to each part
3. Combine the results

Split-Apply-Combine

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

Split-Apply-Combine

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- mutate()
- transmute()

group_by(Sex)

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

Split-Apply-Combine

- filter()
- arrange()
- select()
- rename()
- distinct()
- **group_by()**
- **summarise()**
- mutate()
- transmute()

```
group_by(Sex)
summarise(SexCount = sum(Count))
```

Sex	Morph	Count
Female	Dark	4
Female	Light	6
Male	Dark	8
Male	Light	2

Split-Apply-Combine

- filter()
- arrange()
- select()
- rename()
- distinct()
- **group_by()**
- **summarise()**
- mutate()
- transmute()

```
group_by(Sex)
summarise(SexCount = sum(Count))
```

Sex	SexCount
Male	10
Female	10

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- **mutate()**
- transmute()

dplyr verbs

- filter()
- arrange()
- select()
- rename()
- distinct()
- group_by()
- summarise()
- **mutate()**
- transmute()

```
mutate(Proportion = Count /
       sum(Count))
```

Sex	Morph	Count	Proportion
Female	Dark	4	0.2
Female	Light	6	0.3
Male	Dark	8	0.4
Male	Light	2	0.1

%>%

The pipe.

Chaining operations

- The pipe operator, %>%, enables you to chain together a sequence of commands by piping the result from one step into another.

```
DATA %>%
  select <columns> %>%
  filter <rows> %>%
  arrange <by variable>
```

Chaining operations

```
MOTHS %>%
  filter(Morph == "Dark") %>%
  arrange(desc(Count))
```

Sex	Morph	Count
Male	Dark	8
Female	Dark	4

Chaining operations

```
MOTHS %>%
  filter(Morph == "Dark") %>%
  arrange(desc(Count)) %>%
  mutate(Proportion = Count/sum(Count))
```

Sex	Morph	Count	Proportion
Male	Dark	8	0.67
Female	Dark	4	0.33

Chaining operations

```
MOTHS %>%
  select(Sex, Count) %>%
  group_by(Sex) %>%
  summarise(SexCount = n()) %>%
  mutate(Proportion = SexCount/sum(SexCount))
```

Sex	SexCount	Proportion
Female	10	0.5
Male	10	0.5

Group summaries of individual-level data

```
MOTH_MASS %>%
  group_by(Sex) %>%
  summarise(Mass_mean = mean(MASS),
            Mass_sd = sd(MASS))
```

ID	SEX	MASS
A1	M	3.4
A2	F	6.7
...
A20	M	3.8

Group summaries of individual-level data

```
MOTH_MASS %>%
  group_by(Sex) %>%
  summarise(Mass_mean = mean(MASS),
            Mass_sd = sd(MASS))
```

ID	SEX	MASS
A1	M	3.4
A2	F	6.7
...
A20	M	3.8

Sex	Mass_mean	Mass_sd
F	5.90	0.55
M	3.95	0.29

Join

Merging data frames.

Joining tables: individual-level

ID	SEX	MASS
A1	M	3.4
A2	F	6.7
...
A20	M	3.8

+

ID	MORPH
A1	Light
A2	Dark
...	...
A20	Dark

Joining tables: individual-level

ID	SEX	MASS
A1	M	3.4
A2	F	6.7
...
A20	M	3.8

+

ID	MORPH
A1	Light
A2	Dark
...	...
A20	Dark

Joining tables: individual-level

ID	SEX	MASS	MORPH
A1	M	3.4	Light
A2	F	6.7	Dark
...
A20	M	3.8	Dark

Joining tables: group-level

ID	SEX	MASS	MORPH
A1	M	3.4	Light
A2	F	6.7	Dark
...
A20	M	3.8	Dark

+

SEX	MORPH	LOCATION
M	Light	Top
M	Dark	Bottom
F	Light	Bottom
F	Dark	Top

Joining tables: group-level

ID	SEX	MASS	MORPH
A1	M	3.4	Light
A2	F	6.7	Dark
...
A20	M	3.8	Dark

+

SEX	MORPH	LOCATION
M	Light	Top
M	Dark	Bottom
F	Light	Bottom
F	Dark	Top

Joining tables: group-level

ID	SEX	MASS	MORPH	LOCATION
A1	M	3.4	Light	Top
A2	F	6.7	Dark	Top
...
A20	M	3.8	Dark	Bottom

Practical exercise.

Manipulating data with {dplyr}.

Tidying & manipulating data

- Tidy some 'messy' data
- Interrogate historical data from UK weather stations



Putting it together

- Summing across the whole of the 20th century, find the locations that have the highest ratio of **total sun to total rainfall**...

Putting it together

- Summing across the whole of the 20th century, find the locations that have the highest ratio of **total sun to total rainfall**...

HINTS

- You need to **filter** the range of years,
- Group by** location,
- Summarise** total sun and total rain over this period,
- Mutate** these summary variables into a ratio,
- Select** the variables you are interested in,
- Arrange** your ratio in descending order.

Putting it together

- Summing across the whole of the 20th century, find the locations that have the highest ratio of **total sun to total rainfall**...

```
WEATHER %>%
  filter(Year >= 1900, Year <= 1999) %>%
  group_by(Location) %>%
  summarise(totalrain = sum(Rainfall, na.rm = TRUE),
            totalsun = sum(Sun, na.rm = TRUE)) %>%
  mutate(Sun2Rain = totalsun/totalrain) %>%
  select(Location, Sun2Rain) %>%
  arrange(desc(Sun2Rain))
```

{end}

You can now manipulate data.